

Trax

deep learning 6

Raoul Grouls, 7 juni 2022

Machine learning in Python

Imagine some cars to drive 100x slower

To do machine learning, we would ideally want code to:

- Be as **fast** as possible, especially with matrix multiplications
- Handle the **gradient** for us, in order to learn weights with back propagation
- Have an **ecosystem** for machine learning

In terms of speed, Python seems to be a bad choice.



Possible solutions

- Use Python libraries that have their backend in C/C++ (e.g. PyTorch or TensorFlow)
- Learn a faster language, that also supports autograd
- Use Python libraries with a backend in JAX



GO reference version 0.9.17 Build and Tests on Linux/amd64 passing codecov 58% go report A+ stability unstable



Jax is fast



```
import numpy as np
```

```
def selu(x, alpha=1.67, lambda_=1.05):  
    return lambda_ * np.where(x > 0, x, alpha * np.exp(x) - alpha)
```

```
x = np.arange(1000000)
```

```
%timeit selu(x)
```

✓ 5.7s

7.01 ms \pm 100 μ s per loop (mean \pm std. dev. of 7 runs, 100 loops each)

```
import jax
```

```
import jax.numpy as jnp
```

```
def selu(x, alpha=1.67, lambda_=1.05):  
    return lambda_ * jnp.where(x > 0, x, alpha * jnp.exp(x) - alpha)
```

```
x = jnp.arange(1000000)
```

```
%timeit selu(x).block_until_ready()
```

✓ 1.8s

2.24 ms \pm 105 μ s per loop (mean \pm std. dev. of 7 runs, 100 loops each)

```
selu_jit = jax.jit(selu)
```

```
# Warm up
```

```
selu_jit(x).block_until_ready()
```

```
%timeit selu_jit(x).block_until_ready()
```

✓ 5.6s

689 μ s \pm 2.21 μ s per loop (mean \pm std. dev. of 7 runs, 1,000 loops each)

Jax has an ecosystem



<https://project-awesome.org/n2cholas/awesome-jax>

Libraries


- Neural Network Libraries
 - [Flax](#) - Centered on flexibility and clarity.
 - [Haiku](#) - Focused on simplicity, created by the authors of Sonnet at DeepMind.
 - [Objax](#) - Has an object oriented design similar to PyTorch.
 - [Elegy](#) - A High Level API for Deep Learning in JAX. Supports Flax, Haiku, and Trax.
 - [Trax](#) - "Batteries included" deep learning library focused on providing a simple, high-level API.
 - [Jraph](#) - Lightweight graph neural network library.
 - [Neural Tangents](#) - High-level API for specifying neural networks of both TensorFlow and JAX.
 - [HuggingFace](#) - Ecosystem of pretrained Transformers for a wide range of tasks.
 - [Equinox](#) - Callable PyTrees and filtered JIT/grad transformations => new standard for JAX.
- [NumPyro](#) - Probabilistic programming based on the Pyro library.
- [Chex](#) - Utilities to write and test reliable JAX code.
- [Optax](#) - Gradient processing and optimization library.
- [RLax](#) - Library for implementing reinforcement learning agents.
- [JAX, M.D.](#) - Accelerated, differential molecular dynamics.
- [Coax](#) - Turn RL papers into code, the easy way.
- [Distrax](#) - Reimplementation of TensorFlow Probability, containing probabilistic models and samplers.
- [cvxpylayers](#) - Construct differentiable convex optimization layers.
- [TensorLy](#) - Tensor learning made simple.
- [NetKet](#) - Machine Learning toolbox for Quantum Physics.

Trax

TensorFlow 3.0

- In TensorFlow, there is a lot of baggage that it just has to carry because it's backward compatible.
- Trax is a Google project, intended to improve TensorFlow, breaking backward compatibility.
- The focus is clear code and speed; you can literally read Trax source-code and understand what's going on.

Trax — Deep Learning with Clear Code and Speed



Trax is an end-to-end library for deep learning that focuses on clear code and speed. It is actively used and maintained in the [Google Brain team](#). This

<https://github.com/google/trax>

Downsides

- Trax is fairly recent. You will not find that much examples as for PyTorch / TensorFlow
- Even though there is an ecosystem, it is smaller than the PyTorch / TensorFlow ecosystems
- You will need to dive into the source code; a lot of documentation is to be found there.

Trax is clean

Implementations of Adam

```
def update(self, step, grads, weights, slots, opt_params):
    m, v = slots
    learning_rate = opt_params['learning_rate']
    weight_decay_rate = opt_params['weight_decay_rate']
    b1 = opt_params['b1']
    b2 = opt_params['b2']
    eps = opt_params['eps']
    m = (1 - b1) * grads + b1 * m # First moment estimate.
    v = (1 - b2) * (grads ** 2) + b2 * v # Second moment estimate.
    mhat = m / (1 - b1 ** (step + 1)) # Bias correction.
    vhat = v / (1 - b2 ** (step + 1))
    new_weights = ((1 - weight_decay_rate) * weights - (
        learning_rate * mhat / (jnp.sqrt(vhat) + eps))).astype(weights.dtype)
    return new_weights, (m, v)
```

Trax

```
def _resource_apply_sparse(self, grad, var, indices, apply_state=None):
    var_device, var_dtype = var.device, var.dtype.base_dtype
    coefficients = ((apply_state or {}).get((var_device, var_dtype))
        or self._fallback_apply_state(var_device, var_dtype))

    # m_t = beta1 * m + (1 - beta1) * g_t
    m = self.get_slot(var, 'm')
    m_scaled_g_values = grad * coefficients['one_minus_beta_1_t']
    m_t = tf.compat.v1.assign(m, m * coefficients['beta_1_t'],
        use_locking=self._use_locking)

    with tf.control_dependencies([m_t]):
        m_t = self._resource_scatter_add(m, indices, m_scaled_g_values)

    # v_t = beta2 * v + (1 - beta2) * (g_t * g_t)
    v = self.get_slot(var, 'v')
    v_scaled_g_values = (grad * grad) * coefficients['one_minus_beta_2_t']
    v_t = tf.compat.v1.assign(v, v * coefficients['beta_2_t'],
        use_locking=self._use_locking)

    with tf.control_dependencies([v_t]):
        v_t = self._resource_scatter_add(v, indices, v_scaled_g_values)

    if not self.amsgrad:
        v_sqrt = tf.sqrt(v_t)
        var_update = tf.compat.v1.assign_sub(
            var, coefficients['lr'] * m_t / (v_sqrt + coefficients['epsilon']),
            use_locking=self._use_locking)
        return tf.group(*[var_update, m_t, v_t])
    else:
        v_hat = self.get_slot(var, 'vhat')
        v_hat_t = tf.maximum(v_hat, v_t)
        with tf.control_dependencies([v_hat_t]):
            v_hat_t = tf.compat.v1.assign(
                v_hat, v_hat_t, use_locking=self._use_locking)
        v_hat_sqrt = tf.sqrt(v_hat_t)
        var_update = tf.compat.v1.assign_sub(
            var,
            coefficients['lr'] * m_t / (v_hat_sqrt + coefficients['epsilon']),
            use_locking=self._use_locking)
        return tf.group(*[var_update, m_t, v_t, v_hat_t])
```

Tensorflow

```
def _single_tensor_adam(params: List[Tensor],
    grads: List[Tensor],
    exp_avgs: List[Tensor],
    exp_avg_sqs: List[Tensor],
    max_exp_avg_sqs: List[Tensor],
    state_steps: List[Tensor],
    *,
    amsgrad: bool,
    beta1: float,
    beta2: float,
    lr: float,
    weight_decay: float,
    eps: float,
    maximize: bool):

    for i, param in enumerate(params):

        grad = grads[i] if not maximize else -grads[i]
        exp_avg = exp_avgs[i]
        exp_avg_sq = exp_avg_sqs[i]
        step_t = state_steps[i]
        # update step
        step_t += 1
        step = step_t.item()

        bias_correction1 = 1 - beta1 ** step
        bias_correction2 = 1 - beta2 ** step

        if weight_decay != 0:
            grad = grad.add(param, alpha=weight_decay)

        # Decay the first and second moment running average coefficient
        exp_avg.mul_(beta1).add_(grad, alpha=1 - beta1)
        exp_avg_sq.mul_(beta2).addcmul_(grad, grad.conj(), value=1 - beta2)
        if amsgrad:
            # Maintains the maximum of all 2nd moment running avg. till now
            torch.maximum(max_exp_avg_sqs[i], exp_avg_sq, out=max_exp_avg_sqs[i])
            # Use the max. for normalizing running avg. of gradient
            denom = (max_exp_avg_sqs[i].sqrt() / math.sqrt(bias_correction2)).add_(eps)
        else:
            denom = (exp_avg_sq.sqrt() / math.sqrt(bias_correction2)).add_(eps)

        step_size = lr / bias_correction1
        param.addcddiv_(exp_avg, denom, value=-step_size)

def _multi_tensor_adam(params: List[Tensor],
    grads: List[Tensor],
    exp_avgs: List[Tensor],
    exp_avg_sqs: List[Tensor],
    max_exp_avg_sqs: List[Tensor],
    state_steps: List[Tensor],
    *,
    amsgrad: bool,
    beta1: float,
    beta2: float,
    lr: float,
    weight_decay: float,
    eps: float,
    maximize: bool):

    if len(params) == 0:
        return

    # update steps
    torch._foreach_add_(state_steps, 1)

    if maximize:
        grads = torch._foreach_neg(tuple(grads)) # type: ignore[assignment]

    bias_correction1 = [1 - beta1 ** step.item() for step in state_steps]
    bias_correction2 = [1 - beta2 ** step.item() for step in state_steps]
    if weight_decay != 0:
        torch._foreach_add_(grads, params, alpha=weight_decay)

    torch._foreach_mul_(exp_avgs, beta1)
    torch._foreach_add_(exp_avgs, grads, alpha=1 - beta1)

    torch._foreach_mul_(exp_avg_sqs, beta2)
    torch._foreach_addcmul_(exp_avg_sqs, grads, grads, 1 - beta2)

    if amsgrad:
        # Maintains the maximum of all 2nd moment running avg. till now
        max_exp_avg_sqs = torch._foreach_maximum(max_exp_avg_sqs, exp_avg_sqs) # type:

        # Use the max. for normalizing running avg. of gradient
        max_exp_avg_sq_sqrt = torch._foreach_sqrt(max_exp_avg_sqs)
        bias_correction_sqrt = [math.sqrt(bc) for bc in bias_correction2]
        torch._foreach_div_(max_exp_avg_sq_sqrt, bias_correction_sqrt)
        denom = torch._foreach_add(max_exp_avg_sq_sqrt, eps)
    else:
        exp_avg_sq_sqrt = torch._foreach_sqrt(exp_avg_sqs)
        bias_correction_sqrt = [math.sqrt(bc) for bc in bias_correction2]
        torch._foreach_div_(exp_avg_sq_sqrt, bias_correction_sqrt)
        denom = torch._foreach_add(exp_avg_sq_sqrt, eps)

    step_size = [(lr / bc) * -1 for bc in bias_correction1]
    torch._foreach_addcddiv_(params, exp_avgs, denom, step_size)
```

PyTorch

Trax is readable

Linear / Dense layer

Documentation

```
def forward(self, x):  
    """Executes this layer as part of a forward pass through the model.  
  
    Args:  
        x: Tensor of same shape and dtype as the input signature used to  
           initialize this layer.  
  
    Returns:  
        Tensor of same shape and dtype as the input, except the final dimension  
        is the layer's `n_units` value.  
    """
```

Warning/ Error

```
if self._use_bias:  
    if not isinstance(self.weights, (tuple, list)):  
        raise ValueError(f'Weights should be a (w, b) tuple or list; '  
                          f'instead got: {self.weights}')
```

Linear with bias

```
w, b = self.weights  
return jnp.dot(x, w) + b  # Affine map.
```

Linear without bias

```
else:  
    w = self.weights  
    return jnp.dot(x, w)  # Linear map.
```

Trax is readable

```
def update(self, step, grads, weights, slots, opt_params):
    m, v = slots
    learning_rate = opt_params['learning_rate']
    weight_decay_rate = opt_params['weight_decay_rate']
    b1 = opt_params['b1']
    b2 = opt_params['b2']
    eps = opt_params['eps']
    m = (1 - b1) * grads + b1 * m # First moment estimate.
    v = (1 - b2) * (grads ** 2) + b2 * v # Second moment estimate.
    mhat = m / (1 - b1 ** (step + 1)) # Bias correction.
    vhat = v / (1 - b2 ** (step + 1))
    new_weights = ((1 - weight_decay_rate) * weights - (
        learning_rate * mhat / (jnp.sqrt(vhat) + eps))).astype(weights.dtype)
    return new_weights, (m, v)
```

Algorithm 1: *Adam*, our proposed algorithm for stochastic optimization. See section 2 for details, and for a slightly more efficient (but less clear) order of computation. g_t^2 indicates the elementwise square $g_t \odot g_t$. Good default settings for the tested machine learning problems are $\alpha = 0.001$, $\beta_1 = 0.9$, $\beta_2 = 0.999$ and $\epsilon = 10^{-8}$. All operations on vectors are element-wise. With β_1^t and β_2^t we denote β_1 and β_2 to the power t .

Require: α : Stepsize

Require: $\beta_1, \beta_2 \in [0, 1)$: Exponential decay rates for the moment estimates

Require: $f(\theta)$: Stochastic objective function with parameters θ

Require: θ_0 : Initial parameter vector

$m_0 \leftarrow 0$ (Initialize 1st moment vector)

$v_0 \leftarrow 0$ (Initialize 2nd moment vector)

$t \leftarrow 0$ (Initialize timestep)

while θ_t not converged **do**

$t \leftarrow t + 1$

$g_t \leftarrow \nabla_{\theta} f_t(\theta_{t-1})$ (Get gradients w.r.t. stochastic objective at timestep t)

$m_t \leftarrow \beta_1 \cdot m_{t-1} + (1 - \beta_1) \cdot g_t$ (Update biased first moment estimate)

$v_t \leftarrow \beta_2 \cdot v_{t-1} + (1 - \beta_2) \cdot g_t^2$ (Update biased second raw moment estimate)

$\hat{m}_t \leftarrow m_t / (1 - \beta_1^t)$ (Compute bias-corrected first moment estimate)

$\hat{v}_t \leftarrow v_t / (1 - \beta_2^t)$ (Compute bias-corrected second raw moment estimate)


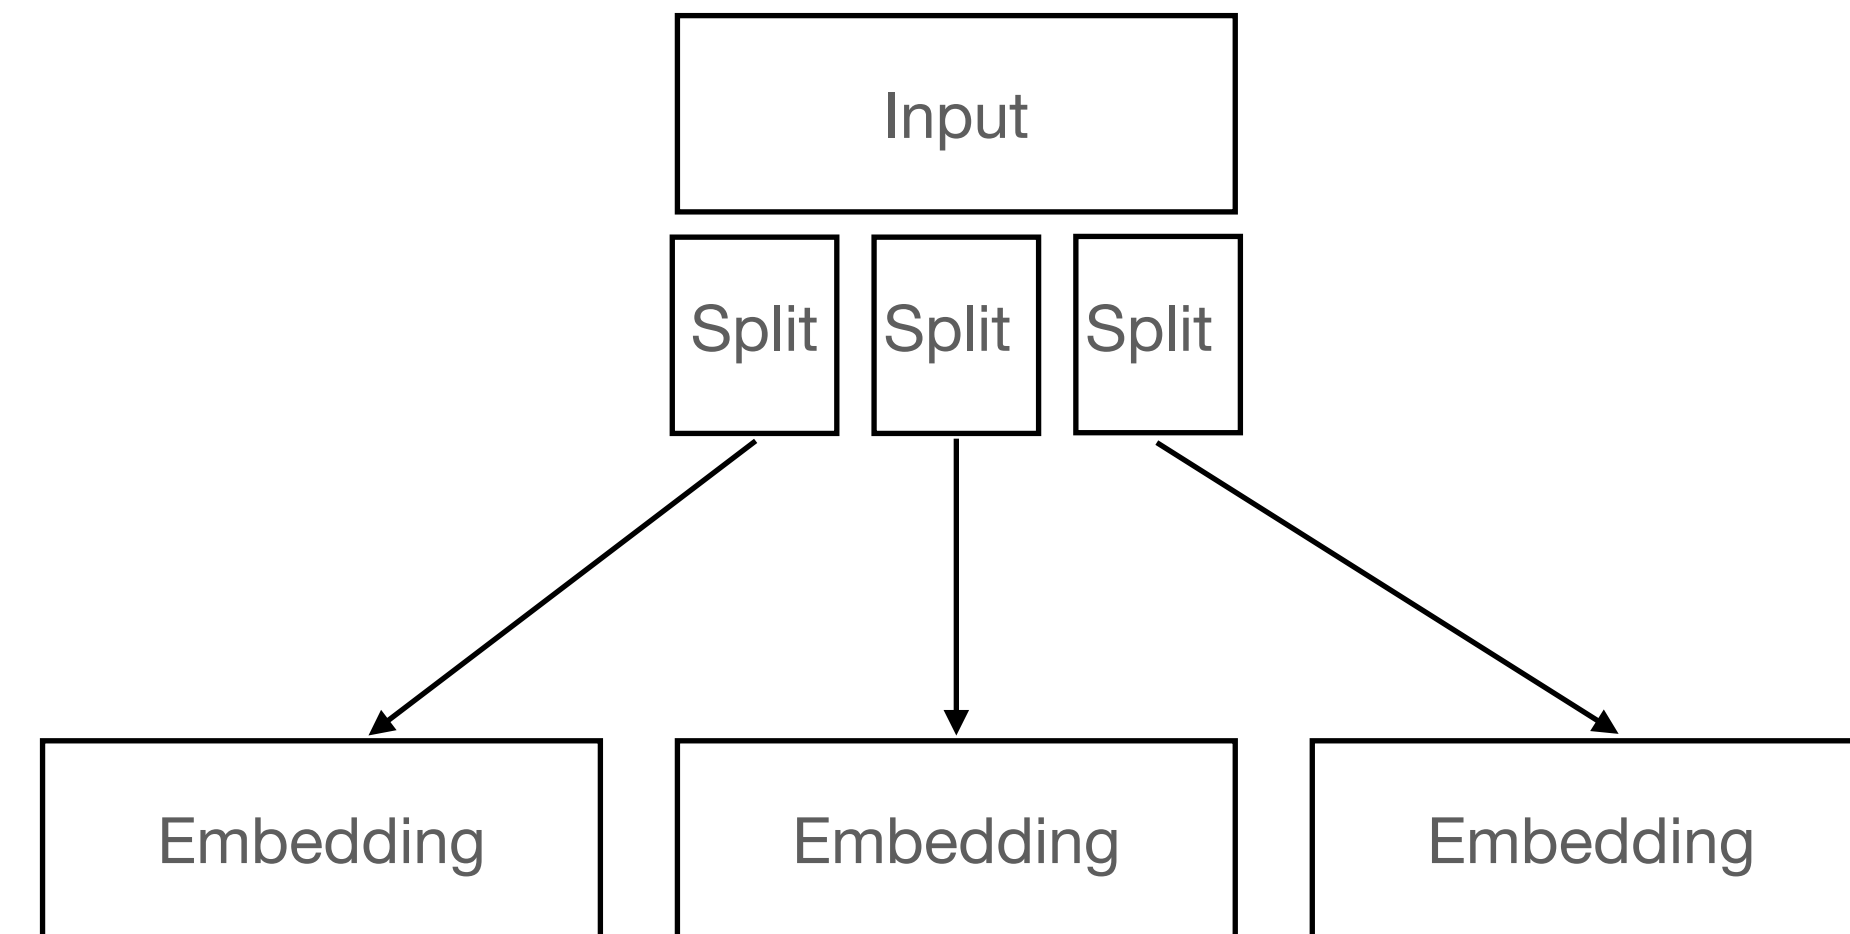
$\theta_t \leftarrow \theta_{t-1} - \alpha \cdot \hat{m}_t / (\sqrt{\hat{v}_t} + \epsilon)$ (Update parameters)

end while

return θ_t (Resulting parameters)

Parallel flows

- In modern architectures, we will often want to implement models with parallel functions
- This can be done with plain Python, but it adds to the complexity
- Trax has multiple “combinators” designed to handle parallel flows



Parallel flows

Multi-embedding

```
emblast = [tl.Embedding(car, d_feature) for car in cardinality]
multiembedding = tl.Parallel(*emblast)
```

Trax implementation with a Parallel layer

Part of the 163 line long implementation in the PyTorch-Forecasting library

```
def init_embeddings(self):
    self.embeddings = nn.ModuleDict()
    for name in self.embedding_sizes.keys():
        embedding_size = self.embedding_sizes[name][1]
        if self.max_embedding_size is not None:
            embedding_size = min(embedding_size, self.max_embedding_size)
        # convert to list to become mutable
        self.embedding_sizes[name] = list(self.embedding_sizes[name])
        self.embedding_sizes[name][1] = embedding_size
    ..
```

```
input_vectors = {}
for name, emb in self.embeddings.items():
    if name in self.categorical_groups:
        input_vectors[name] = emb(
            x[
                ...,
                [self.x_categoricals.index(cat_name) for cat_name in self.categorical_groups[name]]
            ]
        )
    else:
        input_vectors[name] = emb(x[..., self.x_categoricals.index(name)])
```